

AI Bootcamp

Introduction to Natural Language Processing (NLP)

Module 20 Day 1



Class Objectives

By the end of class, you will be able to:

- 1 Define NLP and implement its workflow.
- 2 Demonstrate how to tokenize text.
- 3 Proficiently preprocess text, including tokenization and punctuation handling, for analysis.
- 4 Manage and process punctuation marks and other non-alphabetic characters.
- 5 Differentiate between stemming and lemmatization.
- 6 Understand the importance of removing stopwords.
- 7 Understand and demonstrate how to count tokens and n-grams.



Welcome



Instructor **Demonstration**

Introduction to NLP



What is **Natural Language Processing (NLP)**?





—Jacob **Eisenstein**

Methods for building computer software that understands, generates, and manipulates human language.



What is **NLP** used for?



What is NLP used for?

- Virtual Assistants
- Spam Filters and Fraud Detection
- Translation
- Sentiment Analysis
- Text Summarization
- Content Recommendation
- Speech Recognition
- Healthcare
- Law
- Stock Market
- Text Analytics
- Education

NLP

Most industries have large quantities of textual data that can't be efficiently processed manually.

01

Law: Research, notes, documents, records of legal transactions, governmental information

02

Medical Research: Patient information and history, clinical notes, symptoms

03

Stock Market Analysis: Company disclosures, news articles, report narratives

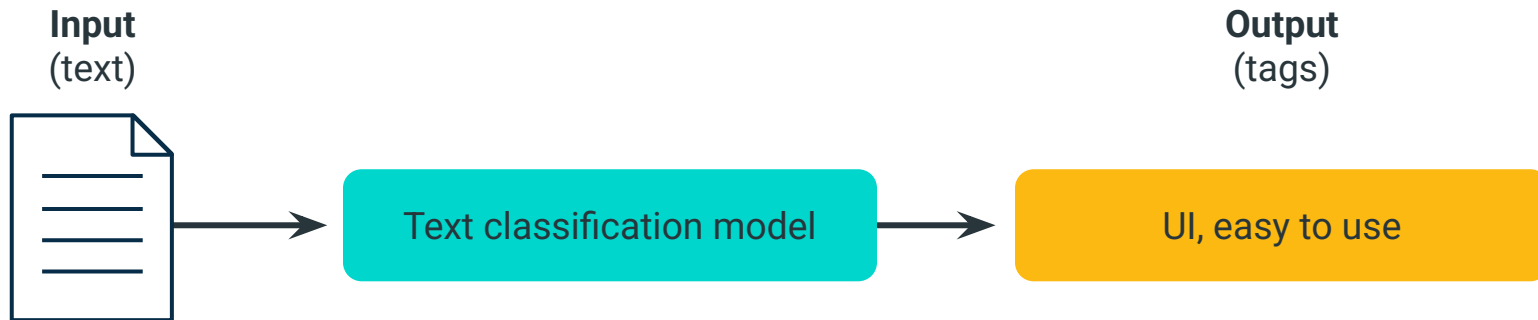


A Few **NLP** Applications




Text Classification

Classifying statements as subjective/objective, positive/negative; finding the reading level or genre of a text



Information Extraction

Finding the diagnosis from a doctor's notes; identifying names of individuals from a witness statement



Applications of Information Extraction

01

Document Summarization

Generating a headline or abstract for a document

02

Complex Question Answering

Answering a question about a subject, given resources or a document on that subject

03

Analyzing Content

Analyzing text or comments to identify trends, such as customer feedback

NLP is Hard

Humans intuitively interpret natural language, but even we aren't great at it all the time. Natural language is:

1


Contextual: The meaning of text depends on situation, speaker, and listener.

2

Ambiguous: Words have multiple meanings and can mean different things in different contexts

3

Nonstandard: There is no general set of rules, especially across dialects, groups, etc.



Natural Languages vs. Computer Languages

Computer languages (programming languages) are:

- Unambiguous
- Based on mathematical logic
- Designed to encode a very specific set of instructions

In order to bridge the gap between human natural language interpretation and processing by a computer, text data must be parsed, organized, and/or encoded. In other words, it must be converted to numbers.

NLP Workflow

01

Preprocessing: Preparing the text, including ingestion.

02

Extraction: Getting interesting features of the text.

03

Analysis: Summarizing these features.

04

Representation: Visualizing your analysis.



Tokenization






What is a **Corpus**?



Corpus

A corpus (plural, corpora) is a large, structured, and organized collection of text documents that usually focus on a specific subject.

A corpus may
contain texts in a
single language
(monolingual corpus)
or text data in
multiple languages
(multilingual corpus).



Tokenization

The process of segmenting running text into words, sentences, or phrases.



Text needs to be segmented into units in order for any processing to be done.



A token is a group of characters that have meaning. It can be words, sentences, or phrases.



Sometimes characters such as punctuation are discarded.



Tokenization is similar to using `split()` in Python.

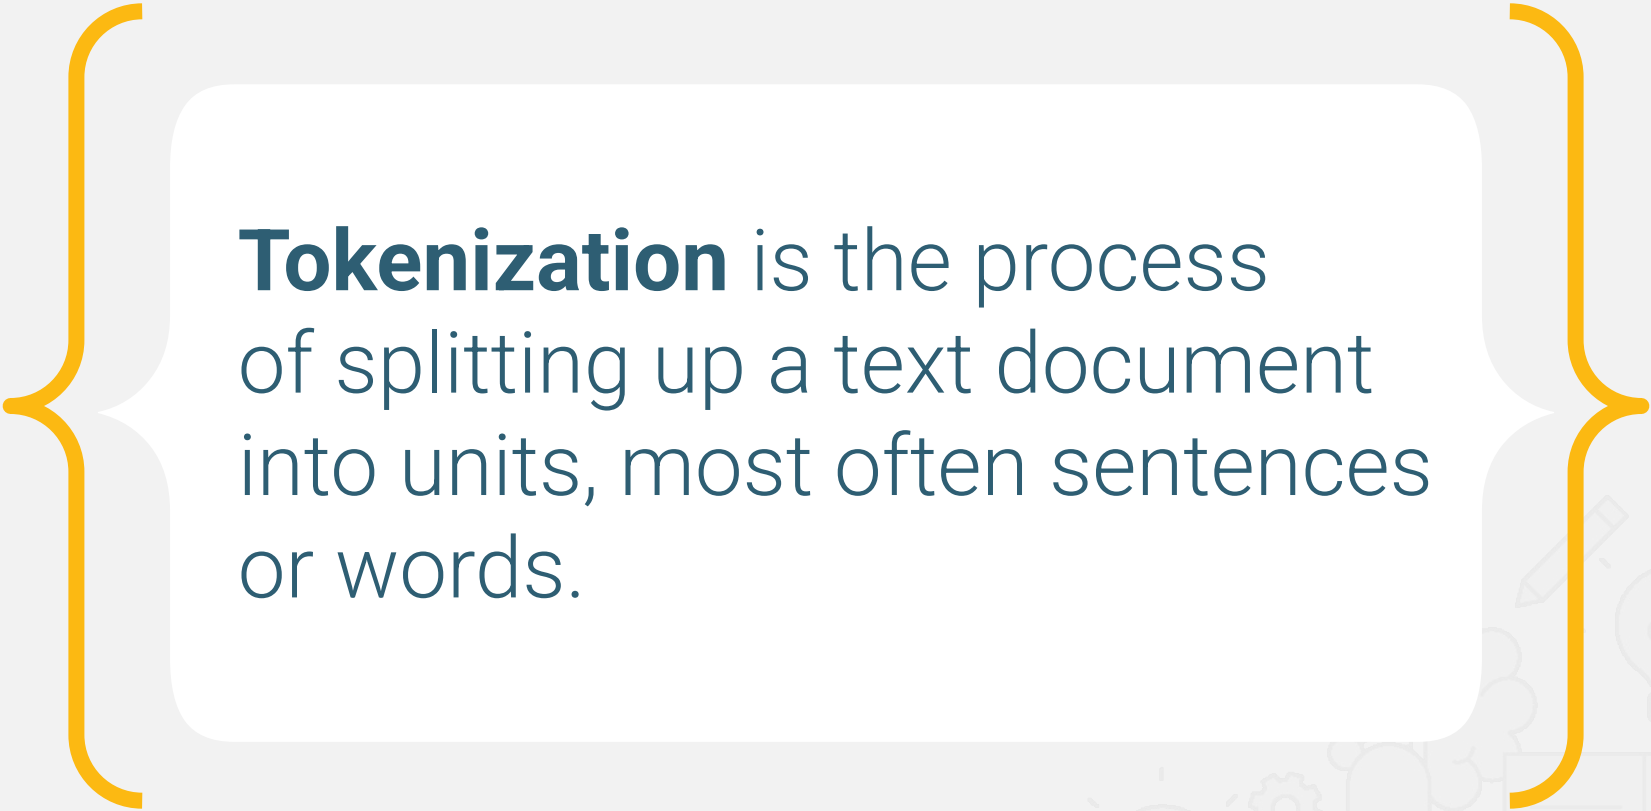


Sentence segmentation and tokenization are often the first steps in an NLP pipeline.

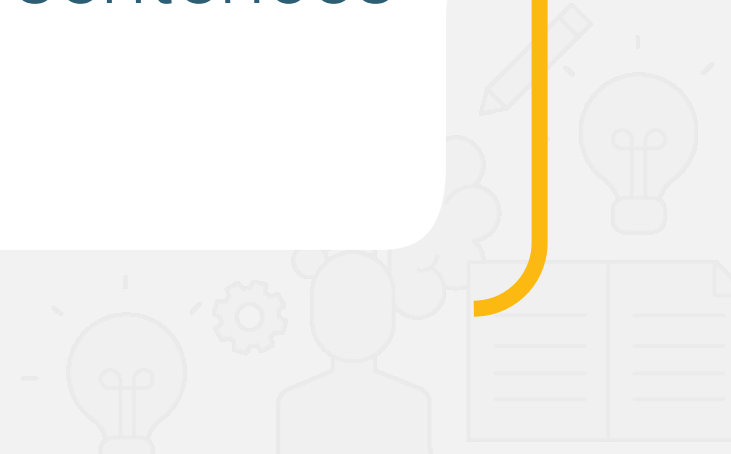
Let's eat, Grandpa!



```
["let's", "eat", "grandpa"]
```



Tokenization is the process of splitting up a text document into units, most often sentences or words.





Instructor **Demonstration**

Tokenization



Activity:

Tokenizing Reuters

In this activity, you will practice both sentence and word tokenization on some articles from the Reuters Corpus.

Suggested Time:

15 Minutes






Time's up!
Let's review



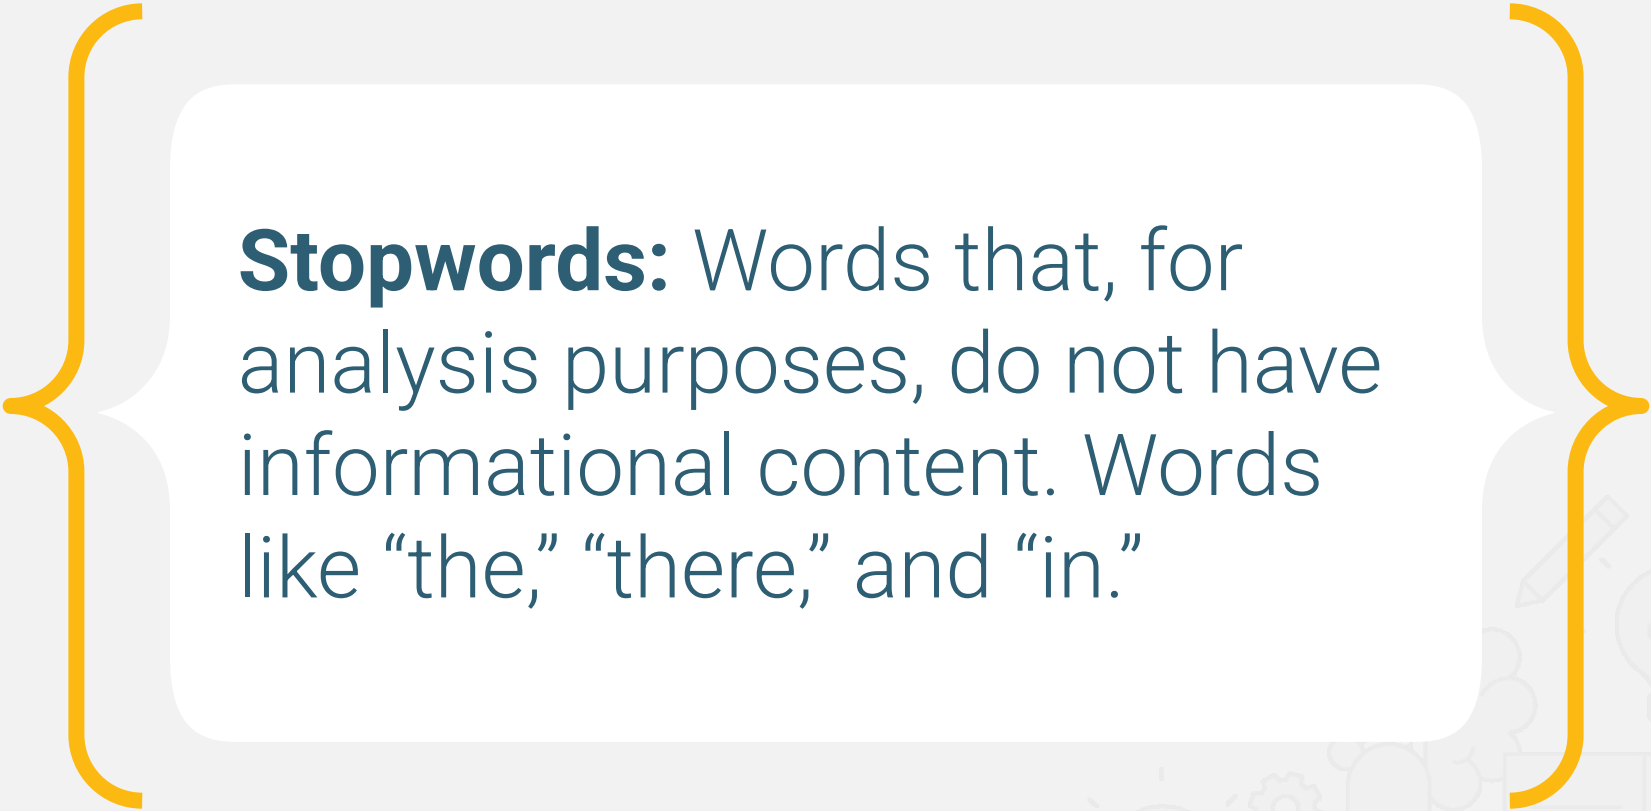
Questions?



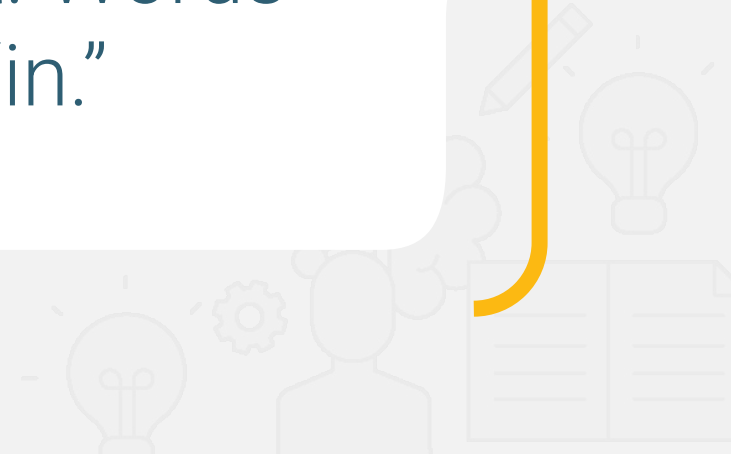


Instructor **Demonstration**

Stopwords



Stopwords: Words that, for analysis purposes, do not have informational content. Words like “the,” “there,” and “in.”



Stopwords

Stopwords are words that are useful for grammar and syntax, but they don't contain any important content.



Generally, stopwords are the most commonly used words in the document.



Examples: *this, to, the, a, there, an*



Stopwords are often removed because they don't distinguish between relevant and irrelevant content.



Activity:

Crude Oil Stopwords

In this activity, you will practice creating a function that performs the preprocessing steps on a news article about crude oil.

Suggested Time:

15 Minutes





Time's up!
Let's review




Questions?





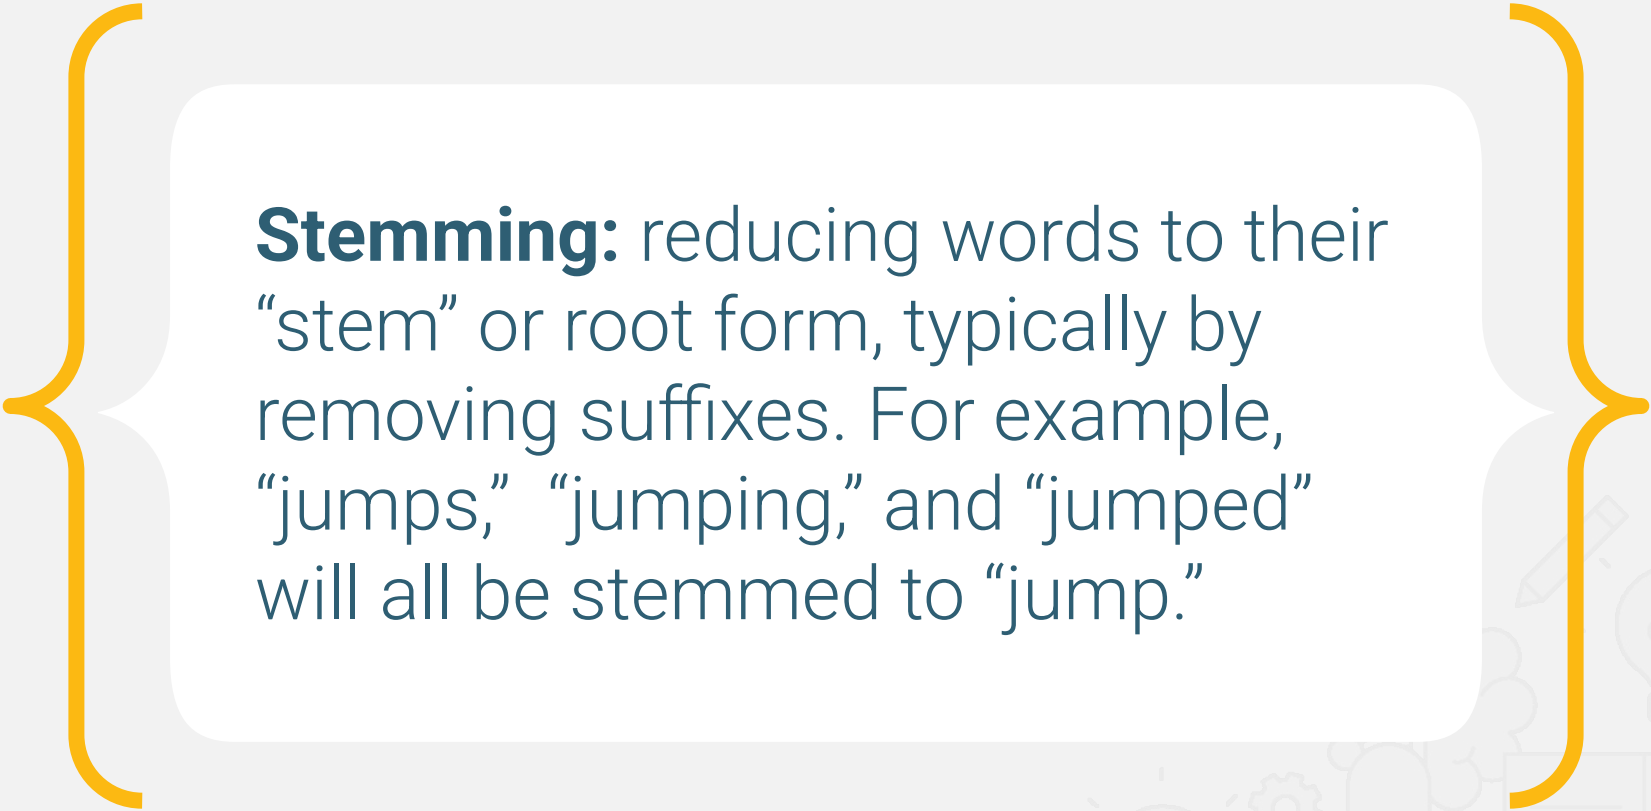
Break

15 mins

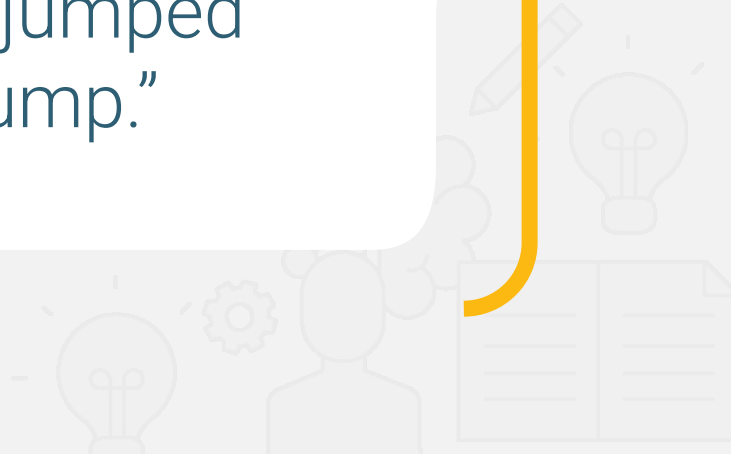


Instructor **Demonstration**

Stemming and Lemmatization




Stemming: reducing words to their “stem” or root form, typically by removing suffixes. For example, “jumps,” “jumping,” and “jumped” will all be stemmed to “jump.”



Stemming

If you were to search the word “boat,” the search results would include results containing the words “boats,” “boater,” and “boating” as well. The word “boat” was the stem for all these words.



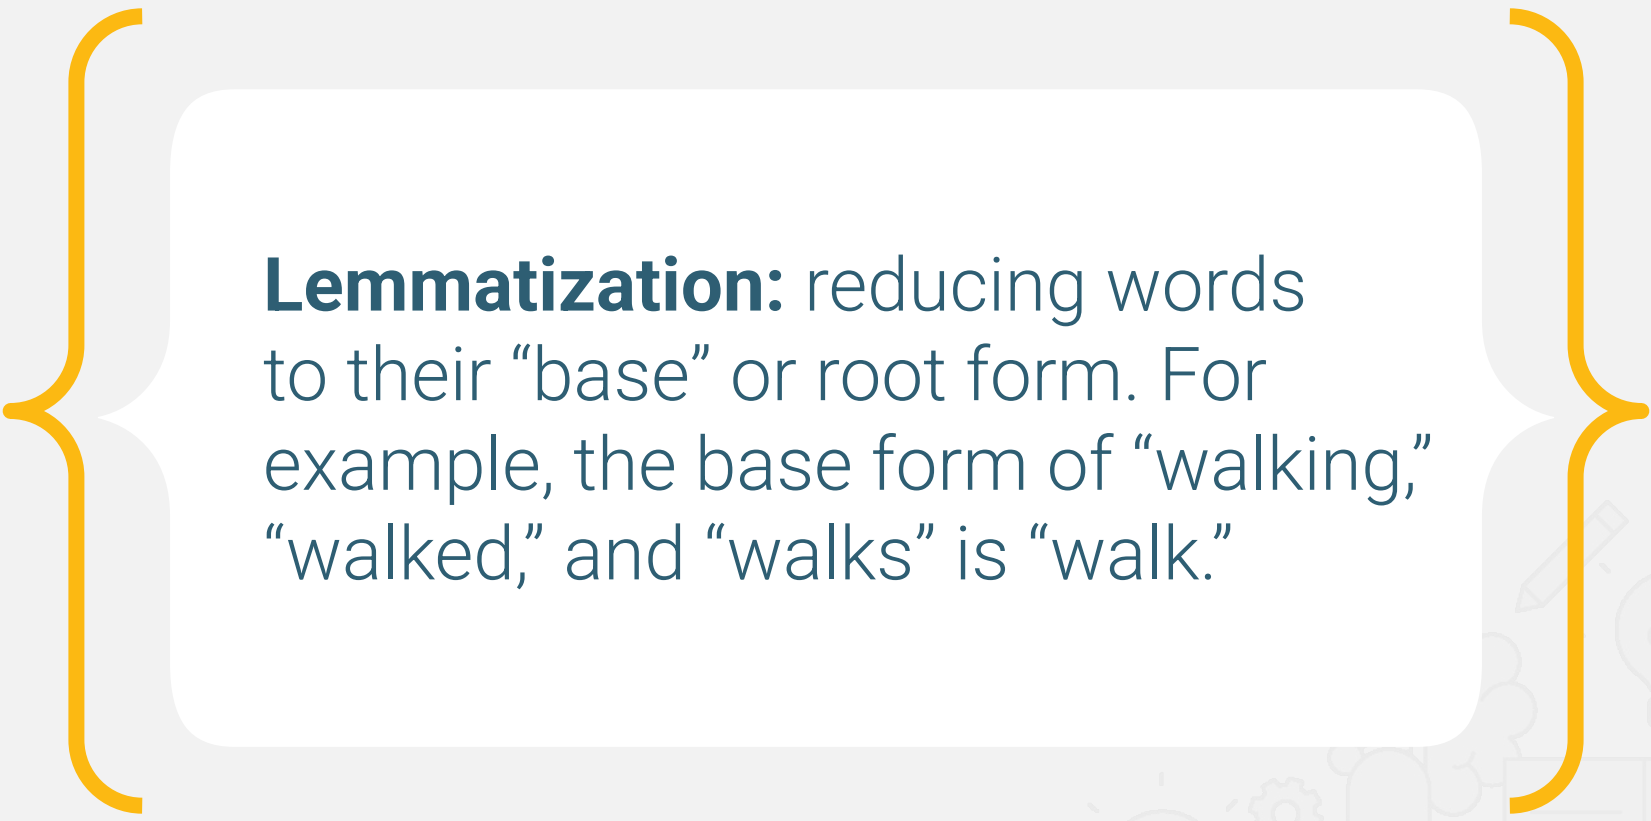
Stemming rules for plural words

SSES $\xrightarrow{\text{maps to}}$ SS

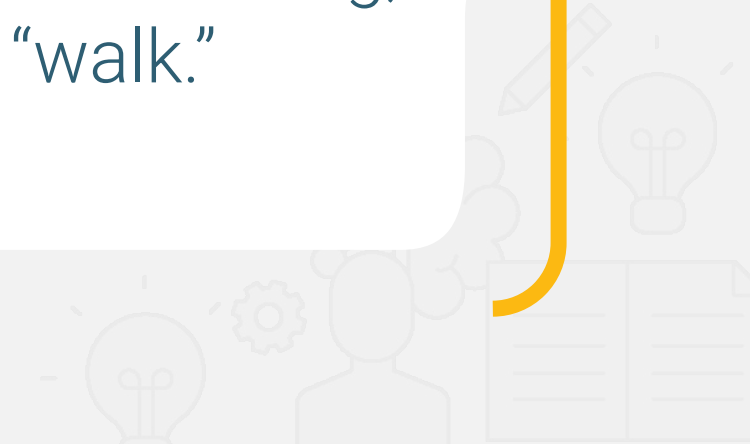
IES $\xrightarrow{\text{maps to}}$ I

SS $\xrightarrow{\text{maps to}}$ SS

S $\xrightarrow{\text{maps to}}$ singular



Lemmatization: reducing words to their “base” or root form. For example, the base form of “walking,” “walked,” and “walks” is “walk.”





Activity:

Stemming and Lemmatization of *Moby Dick*

In this activity, you will write a function that performs the preprocessing steps of removing stopwords and filtering out non-letter characters using regular expressions, as well as applying tokenizing, stemming and lemmatization to the American classic novel *Moby Dick*, written by Herman Melville.

Suggested Time:

15 Minutes




Time's up!
Let's review



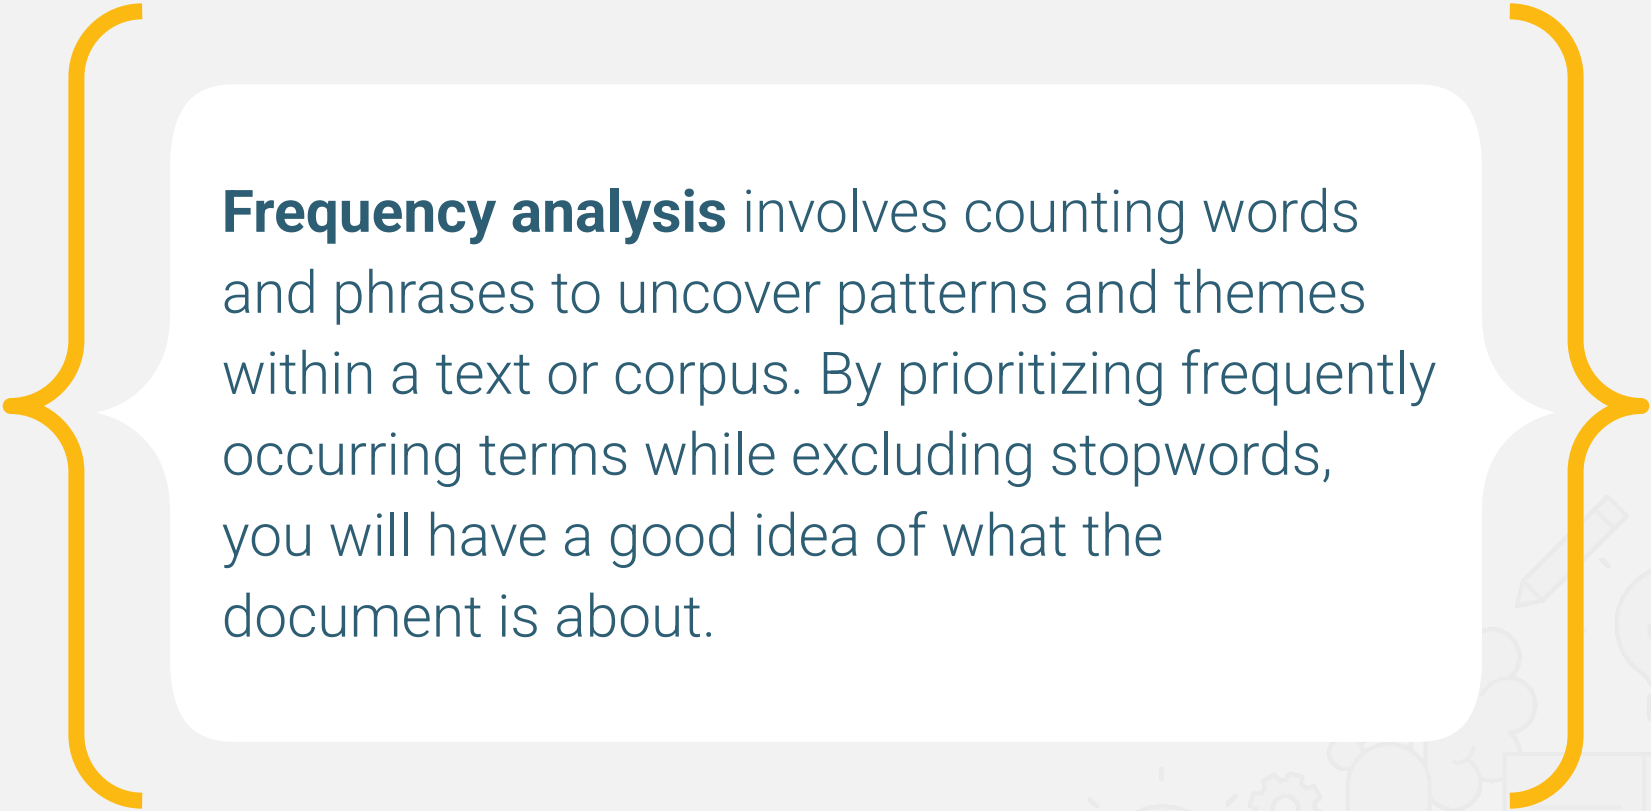
Questions?






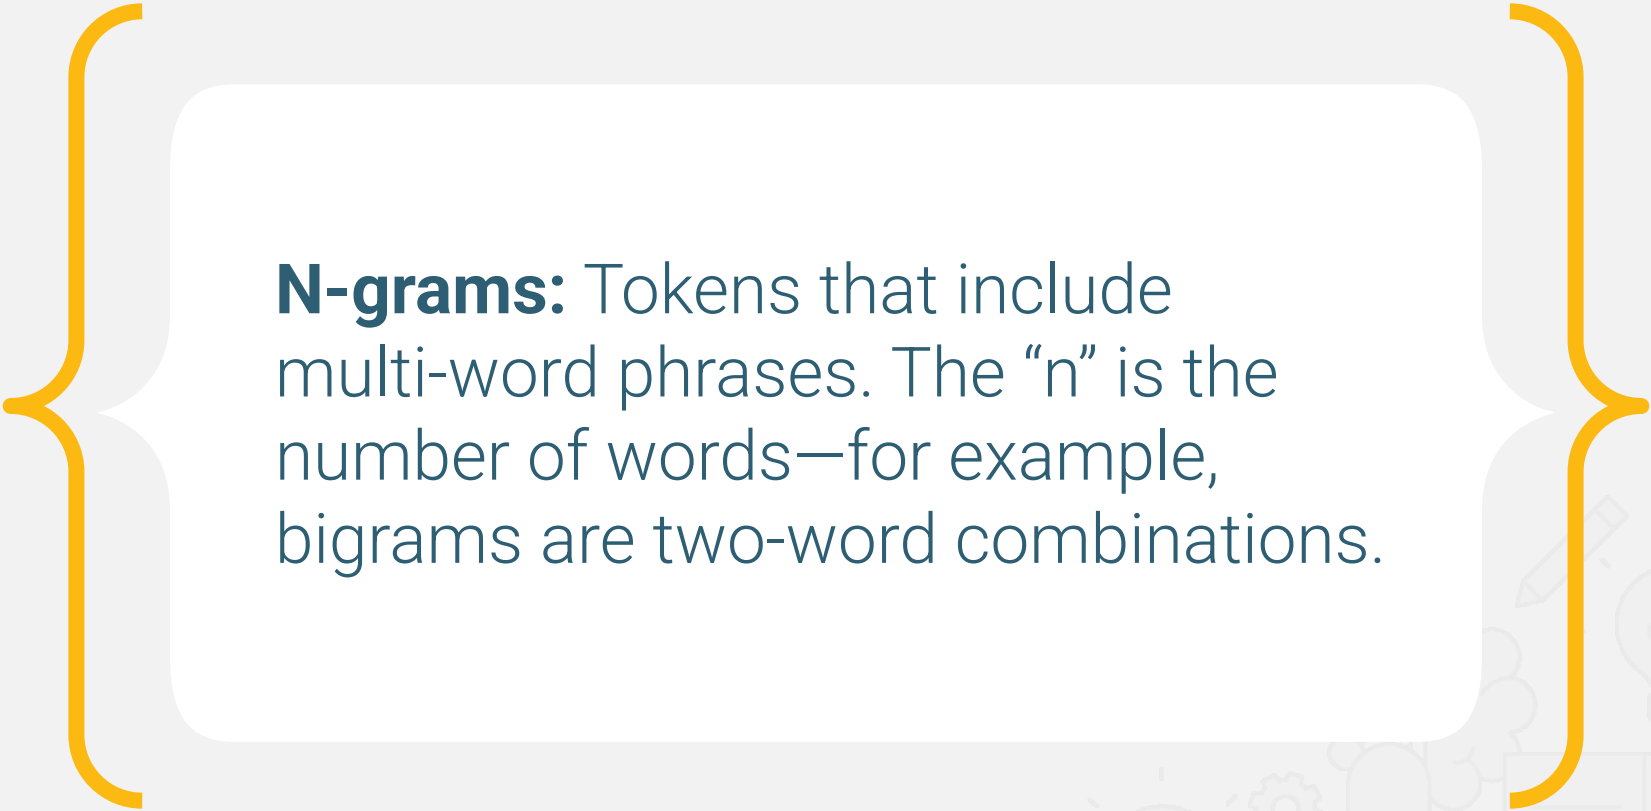
Instructor **Demonstration**

N-gram Counter

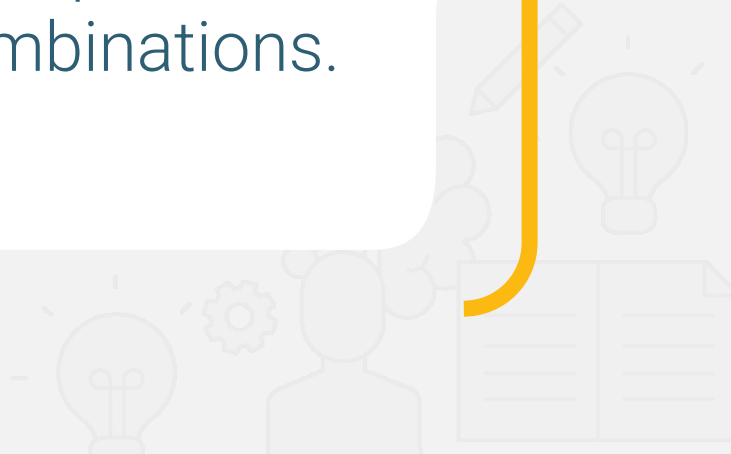


Frequency analysis involves counting words and phrases to uncover patterns and themes within a text or corpus. By prioritizing frequently occurring terms while excluding stopwords, you will have a good idea of what the document is about.





N-grams: Tokens that include multi-word phrases. The “n” is the number of words—for example, bigrams are two-word combinations.



N-grams

A group of n words appearing in sequence from a text.



Splitting on single words can result in a model where syntax and order are ignored.



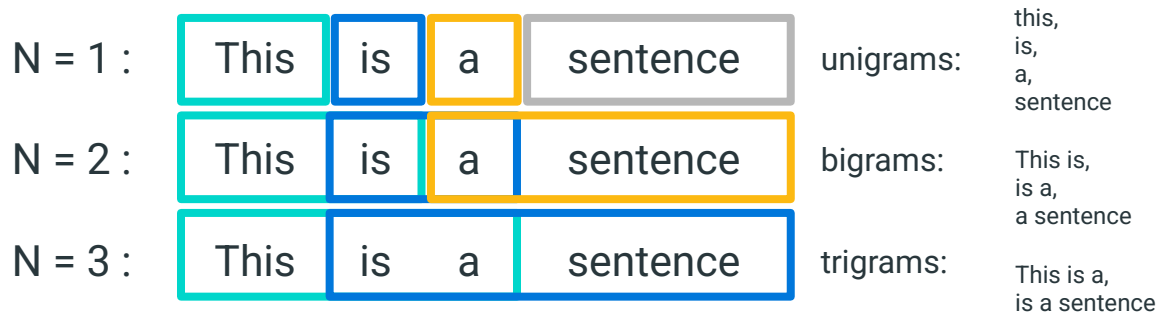
Using an n -gram can be helpful in identifying the multi-word expressions or phrases.



N -grams can be used to calculate how often words follow one another and are applied in generating text (predictive keyboards).



N -grams are helpful in applications like sentiment analysis, where the ordering of the words is important to the context.





Activity:

Word and Bigram Corpus Counter

In this activity, you will create two DataFrames, one that has the top 10 most common words and another that has the top 10 most common bigrams from Reuters articles on grain.

Suggested Time:

15 Minutes






Time's up!
Let's review



Questions?





Let's recap



Recap

After today's lesson you are able to:

- 1 Define NLP and implement its workflow.
- 2 Demonstrate how to tokenize text.
- 3 Proficiently preprocess text, including tokenization and punctuation handling, for analysis.
- 4 Manage and process punctuation marks and other non-alphabetic characters.
- 5 Differentiate between stemming and lemmatization.
- 6 Understand the importance of removing stopwords.
- 7 Understand and demonstrate how to count tokens and n-grams.



Next

In the next lesson, you'll learn advanced NLP techniques such as determining the importance of a word or words in a document, use supervised learning to classify the sentiment of text, and be introduced to the NLP tool spaCy, which has efficient and fast capabilities for tasks like tokenization, part-of-speech tagging, and more.



Questions?





The End